

Name: Aziz Sayyad

Roll No: 381069

Batch: A3

Assignment No: 6

Problem Statement:

Develop a Machine Translation system to translate public information content between English and any Indian language.

Objective:

- To understand the fundamentals of Neural Machine Translation (NMT) and sequence-to-sequence models.
- To load and utilize a pre-trained MarianMT model for English-to-Hindi translation using Hugging Face Transformers.
- To tokenize English input text and generate Hindi translations using autoregressive decoding.
- To translate real-world public information content such as government announcements and health advisories.
- To evaluate the quality and fluency of machine-translated outputs for Indian language content.

S/W Packages and H/W apparatus used:

- Operating System: Windows/Linux/macOS
- Kernel: Python 3.x
- Tools: Jupyter Notebook / Anaconda
- Hardware: CPU with minimum 4GB RAM; GPU recommended for faster inference
- Libraries and packages used: Transformers (Hugging Face), PyTorch, SentencePiece, sacremoses

Theory:

Machine Translation (MT) is the task of automatically converting text from one natural language to another while preserving meaning. It is one of the oldest and most impactful problems in NLP, with modern systems achieving near-human quality on many language pairs using deep learning.

Neural Machine Translation (NMT):

NMT uses encoder-decoder neural networks to translate text end-to-end. The encoder reads and encodes the source sentence into a continuous representation (context vector), and the decoder generates the target sentence word by word conditioned on this representation. Modern NMT systems are based entirely on the Transformer architecture, which uses self-attention and cross-attention mechanisms instead of recurrent networks, enabling better parallelization and capturing long-range dependencies more effectively.

MarianMT:

MarianMT is a family of fast, lightweight Neural Machine Translation models developed by the University of Edinburgh and Helsinki-NLP. They are trained on the OPUS parallel corpus collection, which aggregates millions of translated sentence pairs from sources such as Wikipedia, EU legal

documents, and news. The Helsinki-NLP/opus-mt-en-hi model is specifically fine-tuned for English-to-Hindi translation. MarianMT models use the standard Transformer encoder-decoder architecture with SentencePiece tokenization for subword-level vocabulary sharing.

Autoregressive Decoding:

During translation, the decoder generates output tokens one at a time, with each token conditioned on all previously generated tokens and the encoder's output. The `model.generate()` method automates this process and supports strategies such as greedy decoding, beam search, and sampling. Beam search (the default) maintains multiple candidate translations simultaneously and selects the one with the highest cumulative probability, improving translation quality over simple greedy decoding.

SentencePiece Tokenization:

MarianMT uses SentencePiece, a language-independent subword tokenizer that segments text into subword units (e.g., 'vaccination' -> 'vac', '##cin', '##ation'). This shared vocabulary between source and target languages reduces out-of-vocabulary issues and handles rare or morphologically complex words in both English and Hindi.

Methodology:

- Install required libraries: transformers, torch, and sentencepiece.
- Import AutoTokenizer and AutoModelForSeq2SeqLM from the Hugging Face Transformers library.
- Load the pre-trained Helsinki-NLP/opus-mt-en-hi model and its corresponding tokenizer from the Hugging Face model hub.
- Define a `translate_en_to_hi(text)` function that tokenizes the input English text using the loaded tokenizer with `return_tensors='pt'`.
- Pass the tokenized inputs to `model.generate()` with `max_length=256` to produce output token IDs via beam search decoding.
- Decode the generated token IDs back to a Hindi string using `tokenizer.decode()` with `skip_special_tokens=True`.
- Test the translation system on two real-world public information sentences: a government vaccination announcement and a public health advisory about masks and social distancing.
- Print and compare the original English sentences with their Hindi translations to assess quality.

Use Cases:

- Use Case 1 - Government Public Health Communication: Translating official health advisories, vaccination drive announcements, and COVID-19 guidelines from English to Hindi enables wider reach to non-English-speaking citizens across India, improving public health compliance.
- Use Case 2 - Multilingual E-Governance Portals: Translating government scheme details, legal notices, and citizen service information into Indian languages helps bridge the digital divide and ensures equitable access to public services for rural and semi-urban populations.
- Use Case 3 - Educational Content Localization: Translating English-medium learning materials, textbooks, and online course content into Hindi and other Indian languages supports students learning in their native language, improving comprehension and academic outcomes.

Advantages:

- Pre-trained MarianMT models provide high-quality English-to-Hindi translation without requiring any training data or fine-tuning.
- The Hugging Face pipeline provides a unified, simple API for loading, tokenizing, and generating translations in just a few lines of code.
- SentencePiece tokenization handles the morphological complexity of Hindi script (Devanagari) effectively.
- Beam search decoding produces more fluent and accurate translations compared to greedy decoding.
- The system is easily extendable to other Indian languages by swapping the model (e.g., opus-mt-en-mr for Marathi, opus-mt-en-ta for Tamil).

Limitations:

- The model may produce grammatically imperfect or semantically shifted translations for domain-specific vocabulary (e.g., medical or legal terms).
- MarianMT is a relatively lightweight model; larger models like mBART-50 or NLLB-200 may produce better translations but require significantly more compute.
- Hindi translations of complex or compound English sentences may lose nuance due to structural differences between English and Hindi grammar.
- The model has no mechanism for handling ambiguous words differently based on context (no word sense disambiguation during translation).
- CPU-based inference is slow for long documents; production deployments require GPU acceleration.

Applications:

- Government and public services: Translating official circulars, RTI responses, and public notices for multilingual citizens.
- Healthcare: Translating patient information leaflets, discharge summaries, and medical instructions into regional languages.
- News and media: Automatically translating breaking news articles from English to Hindi for regional news platforms.
- E-commerce: Translating product descriptions and customer reviews to reach Hindi-speaking consumers across India.
- Legal aid: Translating legal documents and court orders into local languages to improve access to justice for non-English-speaking populations.

Conclusion:

This assignment developed a Machine Translation system for English-to-Hindi translation using the Helsinki-NLP/opus-mt-en-hi MarianMT model from Hugging Face Transformers. The system successfully translated real-world public information content -including a government vaccination announcement and a public health advisory -into Hindi using autoregressive beam search decoding. The pre-trained NMT model demonstrated practical translation capability without requiring any custom training data. This system illustrates how modern transformer-based NMT can be rapidly deployed to bridge language barriers in Indian public communication, with straightforward extensibility to other Indian languages through model substitution.